

LAB REPORT 1

TDDE31

Prepared by: Aswathi Karunakaran Beena(aswka169)
Nandana Krishna (nankr143)

Question 1

a. Maximum temperature

```
from pyspark import SparkContext

sc = SparkContext(appName = "exercise 1")
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
lines = temperature_file.map(lambda line: line.split(";"))

# (key, value) = (year,temperature)
year_temperature = lines.map(lambda x: (x[1][0:4], (x[0],float(x[3]))))

#filter
year_temperature = year_temperature.filter(lambda x: int(x[0])>=1950 and int(x[0])<=2014)

#Get maximum temperatures
max_temperatures = year_temperature.reduceByKey(lambda a,b: (a[0], max(a[1], b[1])))
max_temperatures = max_temperatures.sortBy(ascending = False, keyfunc=lambda k: k[1])

# Following code will save the result into /user/ACCOUNT_NAME/BDA/output folder
max_temperatures.saveAsTextFile("BDA/output")
```

Output

```
(u'2011', (u'123340', 32.5))
(u'1982', (u'123250', 33.8))
(u'2002', (u'123250', 33.3))
(u'1986', (u'123250', 33.2))
(u'1991', (u'123250', 32.7))
(u'2006', (u'123250', 32.7))
(u'1999', (u'123250', 32.4))

(u'1996', (u'133260', 30.8))
(u'1985', (u'133260', 29.8))
(u'2009', (u'133250', 31.5))
(u'2012', (u'133250', 31.3))
(u'1964', (u'124020', 31.2))
```

b. Minimum temperature

```
from pyspark import SparkContext

sc = SparkContext(appName = "exercise 1")
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
lines = temperature_file.map(lambda line: line.split(";"))

# (key, value) = (year, temperature)
year_temperature = lines.map(lambda x: (x[1][0:4], (x[0], float(x[3]))))

#filter
year_temperature = year_temperature.filter(lambda x: int(x[0])>=1950 and int(x[0])<=2014)

#Get min
min_temperatures = year_temperature.reduceByKey(lambda a,b: (a[0], min(a[1], b[1])))
min_temperatures = min_temperatures.sortBy(ascending = False, keyfunc=lambda k: k[1])
#print(max_temperatures.collect())

# Following code will save the result into /user/ACCOUNT_NAME/BDA/output folder

min_temperatures.saveAsTextFile("BDA/output")
```

Output

```
(u'1976', (u'140480', -42.2))
(u'1965', (u'140480', -44.0))
(u'1981', (u'140480', -44.0))
(u'1978', (u'140480', -47.7))
(u'2010', (u'140460', -41.7))
(u'2011', (u'140460', -42.0))
(u'2014', (u'140460', -42.5))
(u'2000', (u'140360', -37.6))
(u'1993', (u'140360', -39.0))
(u'2005', (u'140360', -39.4))
(u'1996', (u'140360', -41.7))
(u'1998', (u'140360', -42.7))
```

Question 2

a.Non distinct

```
from pyspark import SparkContext

sc = SparkContext(appName = "exercise 1")
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
lines = temperature_file.map(lambda line: line.split(";"))

# (key, value) = (year,temperature)
year_temperature = lines.map(lambda x: (x[1][0:7], float(x[3])))

#filter by year and temperature reading > 10
year_temperature = year_temperature.filter(lambda x: int(x[0][0:4])>=1950 and int(x[0][0:4])<=2014)
year_temperature = year_temperature.filter(lambda x: int(x[1])>10)

year_temperature = year_temperature.map(lambda x: (x[0], 1))

count = year_temperature.reduceByKey(lambda a,b: a+b)
count = count.sortBy(ascending = False, keyfunc=lambda k: k[1])

#print(count.collect())

# Following code will save the result into /user/ACCOUNT_NAME/BDA/output folder
count.saveAsTextFile("BDA/output")
```

Output

```
(u'2014-07', 144824)
(u'2011-07', 142587)
(u'2010-07', 138997)
(u'2012-07', 130691)
(u'2013-07', 127454)
(u'2009-07', 127430)
(u'2003-07', 125604)
(u'2011-08', 124676)
(u'2002-07', 123771)
(u'2006-08', 123482)
(u'2002-08', 122292)
(u'2009-08', 121596)
(u'2005-07', 121246)
(u'2013-08', 121222)
(u'2006-07', 121122)
```

b. Distinct

```
from pyspark import SparkContext

sc = SparkContext(appName = "exercise 1")
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
lines = temperature_file.map(lambda line: line.split(";"))

# (key, value) = (year, temperature)
year_temperature = lines.map(lambda x: (x[1][0:7], float(x[3])))

# filtering by year
year_temperature = year_temperature.filter(lambda x: int(x[0][0:4])>=1950 and int(x[0][0:4])<=2014)

# count distinct readings from each station
year_temperature = year_temperature.filter(lambda x: int(x[1])>10).distinct()

year_temperature = year_temperature.map(lambda x: (x[0], 1))

count = year_temperature.reduceByKey(lambda a,b: a+b)
count = count.sortBy(ascending = False, keyfunc=lambda k: k[1])

#print(count.collect())

# Following code will save the result into /user/ACCOUNT_NAME/BDA/output folder
count.saveAsTextFile("BDA/output")
```

Output

```
(u'1998-10', 41)
(u'1978-03', 41)
(u'1978-11', 41)
(u'2005-11', 41)
(u'1957-03', 39)
(u'1990-02', 39)
(u'1999-11', 38)
(u'2008-03', 37)
(u'1956-04', 37)
(u'2011-11', 37)
(u'2011-03', 37)
(u'1959-03', 36)
(u'2007-11', 35)
```

Question 3

```
from pyspark import SparkContext

sc = SparkContext(appName = "exercise 1")
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
lines = temperature_file.map(lambda line: line.split(";"))

# (key, value)
year_temperature = lines.map(lambda x: ((x[1][0:7],x[0]) ,float(x[3])))

#filter
year_temperature = year_temperature.filter(lambda x: int(x[0][0][0:4])>=1960 and
int(x[0][0][0:4])<=2014)

sumcount = year_temperature.combineByKey(lambda avg: (avg, 1), lambda x,avg: (x[0] + avg, x[1] + 1),
lambda x,y: (x[0] + y[0], x[1] + y[1]))
avg_temperature = sumcount.map(lambda (k, (sum, count)): (k, sum/count))

# Following code will save the result into /user/ACCOUNT_NAME/BDA/output folder
avg_temperature.saveAsTextFile("BDA/output")
```

Output

```
((u'1970-11', u'115230'), -3.4191666666666668)
((u'1978-06', u'62180'), 14.60700280112046)
((u'1963-11', u'94660'), -0.18111111111111103)
((u'1986-10', u'98040'), 6.6997311827956985)
((u'1985-12', u'124020'), -14.698387096774193)
((u'1960-06', u'117330'), 14.461111111111111)
((u'2012-01', u'148040'), -7.350134408602156)
((u'1992-09', u'105450'), 9.9455555555555552)
((u'1997-08', u'98160'), 20.53442176870748)
((u'2008-01', u'108320'), 1.74260752688172)
((u'2006-11', u'66110'), 6.7884722222222218)
((u'1984-09', u'72080'), 11.082222222222223)
((u'1980-03', u'192710'), -9.583225806451608)
((u'1989-10', u'123070'), 1.6496815286624207)
((u'1988-02', u'132180'), -6.873593073593074)
((u'1998-10', u'62260'), 9.280243572395124)
((u'1983-01', u'108110'), -0.2677419354838712)
((u'1990-02', u'151290'), -0.3294642857142857)
```

Question 4

```
from pyspark import SparkContext

sc = SparkContext(appName = "exercise 1")
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
precipitation_file = sc.textFile("BDA/input/precipitation-readings.csv")
tempslines = temperature_file.map(lambda line: line.split(";"))
precpslines = precipitation_file.map(lambda line: line.split(";"))

#(station number, temperature)
station_temperature = tempslines.map(lambda x: (x[0],float(x[3])))
#((station number, date), precipitation)
daily_precipitation = precpslines.map(lambda x: ((x[0], x[1]),float(x[3])))
#filter

station_temperature = station_temperature.reduceByKey(lambda x,y: x if x>=y else y)
station_temperature = station_temperature.filter(lambda x: int(x[1])>25 and int(x[1])<30)
daily_precipitation = daily_precipitation.reduceByKey(lambda a,b: a + b)
daily_precipitation = daily_precipitation.map(lambda x: (x[0][0], x[1]))
daily_precipitation = daily_precipitation.reduceByKey(lambda x,y: x if x>=y else y)
daily_precipitation = daily_precipitation.filter(lambda x: int(x[1])>100 and int(x[1])<200)

tempsprecps = station_temperature.join(daily_precipitation)
tempsprecps = tempsprecps.map(lambda x: (x[0], x[1], x[3]))

#print(max_temperatures.collect())

# Following code will save the result into /user/ACCOUNT_NAME/BDA/output folder
tempsprecps.saveAsTextFile("BDA/output")
```

Output

No output

Question 5

```
from pyspark import SparkContext

sc = SparkContext(appName = "exercise 1")
# This path is to the file on hdfs
precipitate_file = sc.textFile("BDA/input/precipitation-readings.csv")
ostergotland_file = sc.textFile("BDA/input/stations-Ostergotland.csv")
osterlines = ostergotland_file.map(lambda line: line.split(";"))
precpslines = precipitate_file.map(lambda line: line.split(";"))

#(station number, temperature)
station_oster = osterlines.map(lambda x: (x[0])).collect()

#((station number, date), precipitation)
monthly_precipitation = precpslines.map(lambda x: ((x[0], x[1][0:7]), (float(x[3]))))

#filter
monthly_precipitation = monthly_precipitation.filter(lambda x: x[0][0] in station_oster)
monthly_precipitation = monthly_precipitation.filter(lambda x: int(x[0][1][0:4]) >= 1993 and int(x[0][1][0:4]) <= 2016)

monthly_precipitation = monthly_precipitation.reduceByKey(lambda a,b: a + b)
monthly_precipitation = monthly_precipitation.map(lambda x: (x[0][1][0:7], (x[1], 1)))
count = monthly_precipitation.reduceByKey(lambda a,b: (a[0] + b[0], a[1]+b[1]))
avg_monthly_precipitation = count.map(lambda a: (a[0], a[1][0]/a[1][1]))

#print(max_temperatures.collect())

# Following code will save the result into /user/ACCOUNT_NAME/BDA/output folder
avg_monthly_precipitation.saveAsTextFile("BDA/output")
```

Output

```
u'1996-11', 67.11666666666665)
(u'2008-10', 59.56666666666669)
(u'2014-05', 58.000000000000014)
(u'2004-12', 24.25)
(u'2001-11', 26.383333333333334)
(u'2011-05', 37.85)
(u'1999-07', 29.083333333333334)
```


(u'2010-02', 52.75000000000005)
(u'2013-08', 54.075)
(u'2010-09', 43.08333333333335)
(u'2013-05', 47.92500000000001)
(u'1998-11', 28.96666666666668)
(u'2002-03', 26.93333333333334)
(u'2013-02', 25.52500000000002)
(u'2000-01', 18.61666666666667)